

Justin Fu

justinjf@gmail.com • +1 (734) 972-9418

EDUCATION

University of California, Berkeley, Berkeley, California, USA

Doctor of Philosophy (Ph.D.) in EECS

Aug 2016 – Aug 2021

- Areas: Deep Learning, Reinforcement Learning, Natural Language Processing
- GPA: 3.89/4.00
- Research Advisor: Professor Sergey Levine

University of California, Berkeley, Berkeley, California, USA

Bachelor of Science (B.S.) in EECS

Aug 2011 – May 2015

- Graduated with Highest Honors in EECS.
- Member of HKN (EECS Honor Society)
- GPA: 3.97 / 4.00
- Research Advisor: Professor Pieter Abbeel

WORK EXPERIENCE

Waymo, Mountain View, CA, USA

Research Scientist

Aug 2021 – Current

- Lead developer for Waymax – an open-source, hardware-accelerated simulator for planning and behavior problems in autonomous driving (<https://github.com/waymo-research/waymax>).
- Worked on multi-agent, RL-finetuning of foundation models, achieving an 6% improvement on the WOSAC benchmark over Waymo’s best internal models.
- Working on combining foundation models for autonomous driving with search/optimization to guarantee reliability and safety of system outputs.

DeepMind, London, UK

Research Intern

Sep 2019 – Jan 2020

- Developed algorithms for inferring objectives in multi-agent games using inverse reinforcement learning. This work was accepted for publication in the Journal of Artificial Intelligence Research (JAIR).

Google, Mountain View, California, USA

Student Researcher / Research Intern

Jun 2018 – Aug 2019

- Developed algorithms for language-based human-computer interfaces by inferring objective functions from language commands using inverse reinforcement learning - this work was accepted to ICLR 2019.

PUBLICATIONS AND PREPRINTS

C. Gulino*, **J. Fu***, W. Luo*, G. Tucker*, E. Bronstein, Y. Lu, J. Harb, X. Pan, Y. Wang, X. Chen, J. Co-Reyes, et. al., “Waymax: An Accelerated, Data-Driven Simulator for Large-Scale Autonomous Driving Research,” *Advances in Neural Information Processing Systems – Dataset Track (NeurIPS)*, In Submission. May 2023.

Y. Lu, **J. Fu**, G. Tucker, X. Pan, E. Bronstein, R. Roelofs, B. Sapp, B. White, A. Faust, S. Whiteson, D. Anguelov, S. Levine, “Imitation Is Not Enough: Robustifying Imitation with Reinforcement Learning for Challenging Driving Scenarios,” *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, In Submission. Mar 2023.

C. Snell, M. Yang, **J. Fu**, Y. Su, S. Levine, “Context-aware Language Modeling for Goal-Oriented Dialogue Systems,” *Findings of the Association for Computational Linguistics: NAACL* Jul 2022.

S. Verma, **J. Fu**, M. Yang, S. Levine, “CHAI: A Chatbot AI for Task-oriented Dialogue with Offline Reinforcement Learning,” *North American Chapter of the Association for Computational Linguistics (NAACL)* Jul 2022.

Oral presentation.

J. Fu, A. Tachetti, J. Perolat, Y. Bachrach, “Evaluating Strategic Structures in Multi-Agent Inverse Reinforcement Learning”, *Journal of Artificial Intelligence Research (JAIR)* Aug 2021.

J. Fu, S. Levine, “Offline Model-based Optimization via Normalized Maximum Likelihood Estimation,” *International Conference on Learning Representations (ICLR)* May 2021.

J. Fu*, M. Norouzi*, O. Nachum*, G. Tucker*, Z. Wang, A. Novikov, M. Yang, M. Zhang, Y. Chen, A. Kumar, C. Paduraru, S. Levine, T. Paine*, “Benchmarks for Deep Off-Policy Evaluation”, *International Conference on Learning Representations (ICLR)* May 2021.

D. Ghosh, A. Gupta, A. Reddy, **J. Fu**, C. Devin, B. Eysenbach, S. Levine, “Learning to reach goals without reinforcement learning”, *International Conference on Learning Representations (ICLR)* May 2021.

J. Fu, A. Kumar, O. Nachum, G. Tucker, S. Levine, “D4RL: Datasets for Deep Data-driven Reinforcement Learning”, *arXiv preprint arXiv:2004.07219* Apr 2020.

M. Janner, **J. Fu**, M. Zhang, S. Levine, “When to Trust Your Model: Model-based Policy Optimization”, *Advances in Neural Information Processing Systems (NeurIPS)* Dec 2019.

A. Kumar*, **J. Fu***, G. Tucker, S. Levine, “Stabilizing Off-Policy Q-Learning via Bootstrapping Error Reduction”, *Advances in Neural Information Processing Systems (NeurIPS)* Dec 2019.

J. Fu, A. Korattikara, S. Levine, S. Guadarrama, “From Language to Goals: Inverse Reinforcement Learning for Vision-Based Instruction Following”, *International Conference on Learning Representations (ICLR)* Sep 2019.

J. Fu*, A. Kumar*, M. Soh, S. Levine, “Diagnosing Bottlenecks in Deep Q-learning Algorithms”, *International Conference on Machine Learning (ICML)* Jun 2019.

J. Fu*, A. Singh*, D. Ghosh, L. Wang, S. Levine, “Variational Inverse Control with Events: A General Framework for Data-Driven Reward Definition”, *Advances in Neural Information Processing Systems (NeurIPS)* Dec 2018.

J. Fu, K. Luo, S. Levine, “Learning Robust Rewards with Adversarial Inverse Reinforcement Learning”, *International Conference on Learning Representations (ICLR)* Apr 2018.

J. Fu*, J.D. Co-Reyes*, S. Levine, “ EX^2 : Exploration with Exemplar Models for Deep Reinforcement Learning”, *Advances in Neural Information Processing Systems (NeurIPS)* Dec 2017.

Spotlight presentation.

C. Finn, K. Yu, **J. Fu**, S. Levine, and P. Abbeel, “Generalizing Skills with Semi-Supervised Reinforcement Learning”, *International Conference on Learning Representations (ICLR)* Apr 2017.

J. Fu, S. Levine, and P. Abbeel, “One-Shot Learning of Manipulation Skills with Online Dynamics Adaptation and Neural Network Priors”, *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, Oct 2016.

PROFESSIONAL SERVICE

Co-Organizer, NeurIPS 2021 Workshop on Offline Reinforcement Learning
Co-Organizer, ICML 2019 Workshop on Imitation, Intent, and Interaction (I3)
Reviewer, Advances in Neural Information Processing Systems (NeurIPS) 2018-2022
Reviewer, International Conference on Machine Learning (ICML) 2019, 2020, 2021, 2023
Reviewer, IEEE Conference on Robotics and Automation (ICRA) 2019, 2021
Reviewer, International Conference on Learning Representations (ICLR) 2019, 2020, 2024

TEACHING EXPERIENCE

CS188 (Introduction to Artificial Intelligence)	
UC Berkeley	Jan 2018 – May 2018
Teaching Assistant	Jun 2020 – Aug 2020
CS221 (Artificial Intelligence: Principles & Techniques)	
Stanford University	
Teaching Assistant	Sep 2015 – Dec 2015
CS61C (Great Ideas in Computer Architecture)	
University of California, Berkeley	
Teaching Assistant	Jun 2013 – Aug 2013

OUTREACH

BAIR AI4ALL Summer Camp
Co-organizer
Project Mentor

Aug 2018
Jul 2017